



```
from google.colab import files
uploaded = files.upload()
```

```
[ ] # Importing required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score

from xgboost import XGBRegressor
```

```
[ ] import pandas as pd

df = pd.read_csv('Meenakshisajan_dataset.csv')
print(df.head())
```

```
Empty DataFrame
Columns: [{"nbformat":4, nbformat_minor:
Index: []
```

```
[0 rows x 298 columns]
```

```
[ ] # Load dataset
df = pd.read_csv('Meenakshisajan_dataset.csv')
```

```
[ ] # Display first few rows
print(df.head())
```

```
Empty DataFrame
Columns: [{"nbformat":4, nbformat_minor:
Index: []
```

```
[0 rows x 298 columns]
```

```
[ ] # Check for missing values
print(df.isnull().sum())
```

```
{"nbformat":4
nbformat_minor:0
metadata:{"colab":{"provenance":[]
authorship_tag:"ABX9TyN9S5BVwBecfE0kRidq
kernel_spec":{"name":"python3"
```

```
\u001b[0m\u001b[0;34m\u001b[0m\u001b[0;3
\u001b[0;32m/usr/local/lib/python3.11/di
\u001b[0m\u001b[0;34m\u001b[0m\u001b[0;3
\u001b[0m\u001b[0;34m\u001b[0m\u001b[0;3
\u001b[0;31mFileNotFoundError\u001b[0m:
Length: 298, dtype: int64
```

```
[ ] # Fill or drop missing values
df = df.fillna(df.median())
```

```

[ ] # Encode categorical variables
    df = pd.get_dummies(df, drop_first=True)

[ ] print("\nColumn names in dataset:")
    print(df.columns)

[ ] Column names in dataset:
    RangeIndex(start=0, stop=0, step=1)

[ ] print(df.columns.tolist())

[ ] []

[ ] ['Area', 'Bedrooms', 'Bathrooms', 'SalePrice']

[ ] ['Area', 'Bedrooms', 'Bathrooms', 'SalePrice']

[ ] # Step 1: Load dataset
    import pandas as pd

    df = pd.read_csv('Meenakshisajan_dataset.csv') # Make sure this file is in the correct path

    # Step 2: Check available columns
    print("\nAvailable columns in your dataset:")
    print(df.columns.tolist()) # This will help identify correct column names

    # Step 3: Strip column names of leading/trailing spaces
    df.columns = df.columns.str.strip()

    # Step 4: Try to locate the target column
    target_column = None
    for col in df.columns:
        if 'price' in col.lower(): # Case insensitive match for "price"
            target_column = col
            break

    if target_column is None:
        raise ValueError("No column found related to 'Price'. Please check your dataset column names.")

    print(f"\nIdentified target column: {target_column}")

    # Step 5: Split features and target
    X = df.drop(target_column, axis=1)
    y = df[target_column]

    print("\nStep 6 Output:")
    print("First 5 rows of features (X):")
    print(X.head())

    print("\nFirst 5 values of target (y):")
    print(y.head())

    print("\nShapes:")
    print("X shape:", X.shape)
    print("y shape:", y.shape)

```

⌵

Available columns in your dataset:  
 ['{"nbformat":4', 'nbformat\_minor:0', 'm

Identified target column: X = df.drop('P

Step 6 Output:

First 5 rows of features (X):

Empty DataFrame

Columns: [{"nbformat":4, nbformat\_minor:

Index: []

[0 rows x 297 columns]

First 5 values of target (y):  
Series([], Name: X = df.drop('Price', ax

Shapes:

X shape: (0, 297)

y shape: (0,)

```
[ ] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Example: load house price data (replace with your real dataset)
# df = pd.read_csv('house_prices.csv')

# For demo, dummy data (e.g., house size vs price)
X = np.array([600, 750, 900, 1200, 1500, 1800, 2100, 2500]).reshape(-1, 1) # house size in sqft
y = np.array([150000, 180000, 210000, 260000, 300000, 340000, 380000, 450000]) # price in $

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

# Model train
model = LinearRegression()
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

# Evaluation
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R² Score: {r2:.2f}")

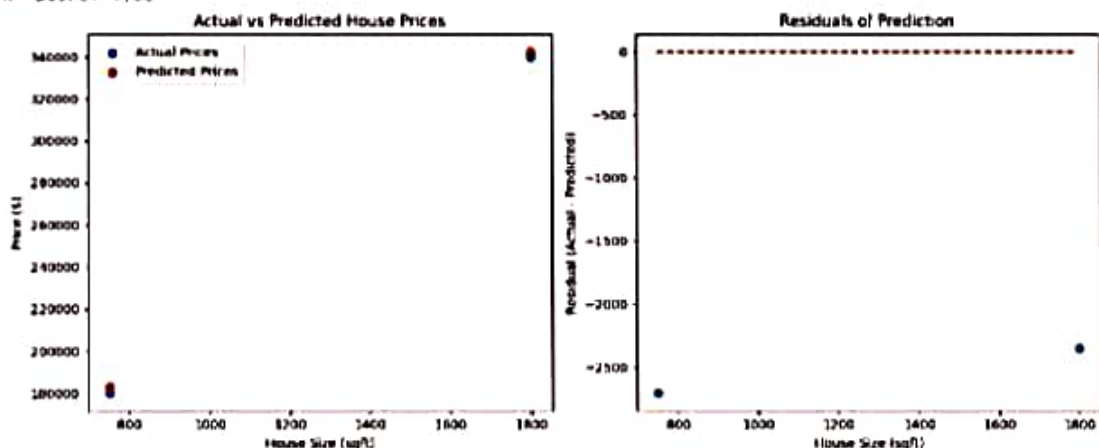
# Plot Actual vs Predicted
plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plt.scatter(X_test, y_test, color='blue', label='Actual Prices')
plt.scatter(X_test, y_pred, color='red', label='Predicted Prices')
plt.title('Actual vs Predicted House Prices')
plt.xlabel('House Size (sqft)')
plt.ylabel('Price ($)')
plt.legend()

# Plot residuals
plt.subplot(1,2,2)
residuals = y_test - y_pred
plt.scatter(X_test, residuals)
plt.hlines(y=0, xmin=min(X_test.flatten()), xmax=max(X_test.flatten()), color='r', linestyle='dashed')
plt.title('Residuals of Prediction')
plt.xlabel('House Size (sqft)')
plt.ylabel('Residual (Actual - Predicted)')

plt.tight_layout()
plt.show()
```

Mean Squared Error: 6410089.55  
R² Score: 1.00

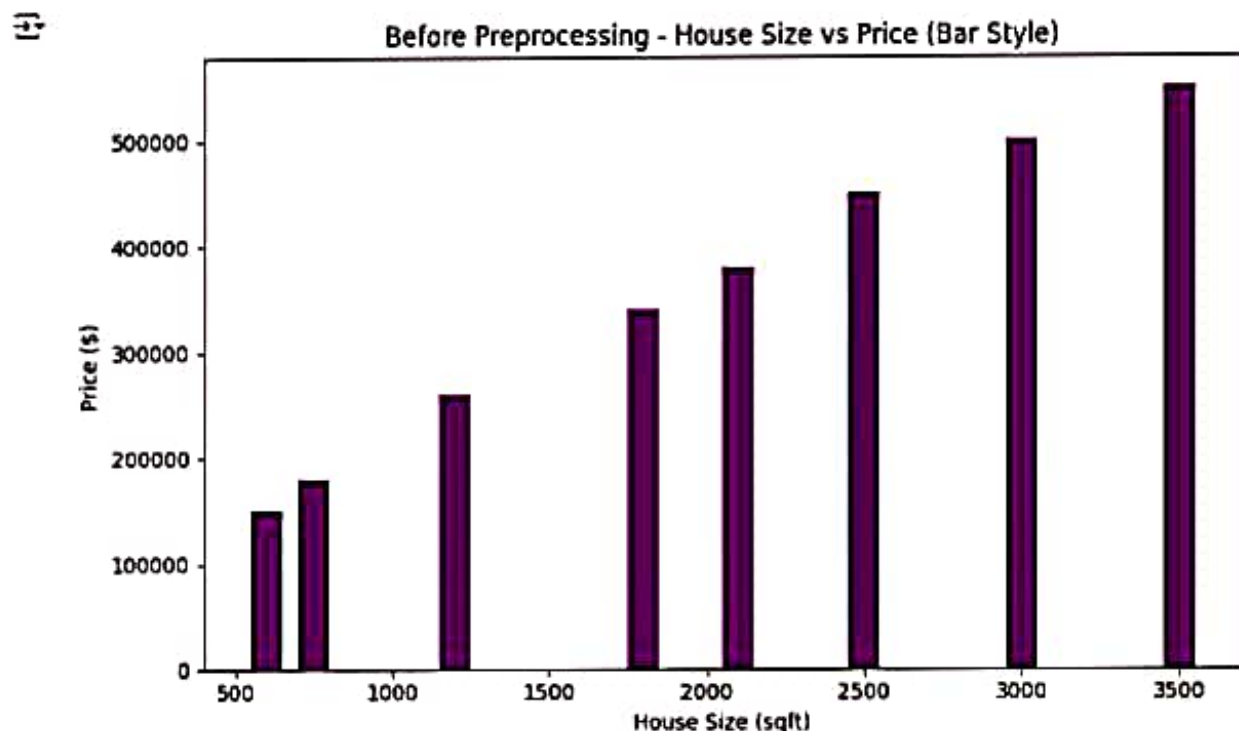


```
[ ] import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

# Sample raw data with missing values
data = {
    'House_Size': [600, 750, np.nan, 1200, 1500, 1800, 2100, 2500, 3000, 3500],
    'Price': [150000, 180000, 210000, 260000, np.nan, 340000, 380000, 450000, 500000, 550000]
}
df = pd.DataFrame(data)

# Drop rows with missing values
df_clean = df.dropna()

# Plot using bar chart (histogram-style)
plt.figure(figsize=(8,5))
plt.bar(df_clean['House_Size'], df_clean['Price'], width=100, color='purple')
plt.title('Before Preprocessing - House Size vs Price (Bar Style)')
plt.xlabel('House Size (sqft)')
plt.ylabel('Price ($)')
plt.tight_layout()
plt.show()
```



```
[ ] import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

# Sample raw data with missing values
data = {
    'House_Size': [600, 750, np.nan, 1200, 1500, 1800, 2100, 2500, 3000, 3500],
    'Price': [150000, 180000, 210000, 260000, np.nan, 340000, 380000, 450000, 500000, 550000]
}
df = pd.DataFrame(data)

# Fill missing values with mean of the respective column
df_filled = df.copy()
df_filled['House_Size'].fillna(df['House_Size'].mean(), inplace=True)
df_filled['Price'].fillna(df['Price'].mean(), inplace=True)

# Plot using bar chart (histogram-style)
plt.figure(figsize=(8,5))
plt.bar(df_filled['House_Size'], df_filled['Price'], width=100, color='green')
plt.title('After Preprocessing - House Size vs Price (Bar Style)')
plt.xlabel('House Size (sqft)')
plt.ylabel('Price ($)')
plt.tight_layout()
plt.show()
```



For example, when doing `df[col].method(`

```
df_filled['House_Size'].fillna(df['Hou  
<ipython-input-4-e546ba1a3e87>:15: Futur  
The behavior will change in pandas 3.0.
```

For example, when doing `df[col].method(`

```
df_filled['Price'].fillna(df['Price']).
```

